

Revealing the Power of Post-Training for Small Language Models via Knowledge Distillation

Miao Rang[†] Zhenni Bi[†] Hang Zhou Hanting Chen An Xiao Tianyu Guo Kai Han^{*} Xinghao Chen[#]
Yunhe Wang[#]

Huawei Noah’s Ark Lab

Abstract

The rapid advancement of large language models (LLMs) has significantly advanced the capabilities of artificial intelligence across various domains. However, their massive scale and high computational costs render them unsuitable for direct deployment in resource-constrained edge environments. This creates a critical need for high-performance small models that can operate efficiently at the edge. Yet, after pre-training alone, these smaller models often fail to meet the performance requirements of complex tasks. To bridge this gap, we introduce a systematic post-training pipeline that efficiently enhances small model accuracy. Our post training pipeline consists of curriculum-based supervised fine-tuning (SFT) and offline on-policy knowledge distillation. The resulting instruction-tuned model achieves state-of-the-art performance among billion-parameter models, demonstrating strong generalization under strict hardware constraints while maintaining competitive accuracy across a variety of tasks. This work provides a practical and efficient solution for developing high-performance language models on Ascend edge devices.

1. Introduction

Large Language Models (LLMs) have transformed the landscape of artificial intelligence by leveraging self-attention mechanisms and massive-scale pre-training to capture hierarchical linguistic patterns, semantic relationships, and cross-domain knowledge (Achiam et al., 2023). Open-sourced models such as LLaMA (Touvron et al., 2023), DeepSeek (Liu et al., 2024a), Qwen (Bai et al., 2023; Yang et al., 2025) and openPangu (Chen et al., 2025; Tang et al.,

2025) have demonstrated exceptional performance in complex tasks including text generation, reasoning, and multi-lingual understanding. However, their success hinges on enormous parameter counts (*e.g.*, DeepSeek-V3 (Liu et al., 2024a) has 671B parameters in total) and extensive computational resources, which limit their deployment in latency-sensitive or resource-constrained environments. The prohibitive energy consumption and hardware requirements of LLMs have raised critical challenges for real-world applications, particularly in scenarios demanding on-device processing, privacy preservation, or low-latency responsiveness.

To address these limitations, a paradigm shift has emerged toward developing efficient Small Language Models (SLMs) for edge devices and resource-constrained platforms. The training pipeline for small models is divided into pre-training and post-training. As is well known, pre-training focuses on building foundational capabilities and requires substantial data and training resources. In contrast, post-training concentrates on enhancing these abilities and typically demands significantly fewer resources. Therefore, leveraging post-training to substantially improve the accuracy of small models under resource constraints has become a critical area of research. Existing post-training methods include supervised fine-tuning (Lobo et al., 2024; Luong et al., 2024), and knowledge distillation (*e.g.*, GKD (Agarwal et al., 2024)).

Building upon these established techniques, our work introduces a multi-stage post-training pipeline designed to enhance small language models(*e.g.*, openPangu Embedded-1B¹ which is specifically developed for efficient inference on Ascend edge devices). Our process begins with SFT using a curated curriculum that transitions from step-by-step reasoning to fast-response examples, thereby strengthening the model’s fundamental instruction-following abilities. The instruct model after SFT is named as openPangu Embedded-1B-SFT.

[†]Equal contribution. ^{*}Project Leader. [#]Corresponding Author.
Correspondence to:

Xinghao Chen <xinghao.chen@huawei.com>,
Yunhe Wang <yunhe.wang@huawei.com>.

¹<https://ai.gitcode.com/ascend-tribe/openPangu-Embedded-1B-V1.1>

Subsequently, we employ offline on-policy knowledge distillation from a larger teacher model that shares the same tokenizer, enabling efficient logit-level knowledge transfer. The resulting instruction-tuned model, named openPangu Embedded-1B-KD, benefits from this comprehensive post-training pipeline and achieves state-of-the-art accuracy among billion-scale instruction models.

2. Related Work

Small Language Models. The development of small language models in the 1 billion parameter range has progressed through several distinct phases, evolving from initial capability demonstrations to highly optimized systems. Initial explorations, such as OPT-1.3 B (Zhang et al., 2022) and GPT-Neo-1.3 B (Kashyap et al., 2022), first established this scale as an efficient sweet spot for natural language understanding (NLU) tasks. Building on this foundation, subsequent research shifted its focus toward the critical role of data quality, with models like DeepSeek-Coder-1.3 B (Guo et al., 2024) demonstrating that curated, domain-specific corpora could enable smaller models to surpass much larger baselines on expert tasks. This data-centric paradigm is best exemplified by Phi-1/1.5 (Fu et al., 2024), which achieve remarkable performance on HumanEval using a small, “textbook-quality” dataset, achieving over 50 % pass@1. The focus then shift to architecture-algorithm co-design in 2024, with innovations like MobileLLM (Liu et al., 2024b) introducing latency-efficient architectures, TinyLlama-1.1 B (Zhang et al., 2024) showcasing the benefits of extreme-scale training on 3 trillion tokens, and the fully open-source OLMo-1 B (Groeneveld et al., 2024) becoming a standard testbed for scaling-law research.

Model-aware Training. Earlier work in LLM post-training primarily focuses on selecting high-quality training data from a general perspective without considering model-specific issues (Liu et al.; Zhou et al., 2023a). While this approach has been effective in improving model performance, recent studies have shown that data distributions significantly deviating from the base model’s can be difficult for the model to learn from and may even degrade performance (Ren et al., 2024). As a result, more research advocates for model-specific data selection (Du et al., 2023; Li et al., 2024). These studies do not explore distillation scenarios, and our work further extends this concept by integrating an iterative distillation pipeline, leveraging our proposed model-aware complexity score.

Knowledge Distillation. Knowledge Distillation (KD) is a widely adopted model compression technique where a compact student model is trained to replicate the behavior of a larger teacher model (Hinton et al., 2015), facilitating the deployment of high-performance models in resource-constrained settings. A fundamental challenge in

applying KD to autoregressive models is the exposure bias phenomenon (Ross et al., 2011)—a discrepancy between the training distribution, which is conditioned on ground-truth sequences, and the inference distribution, where the model conditions on its own generated outputs. This mismatch can lead to an accumulation of errors during generation. To directly mitigate this issue, on-policy distillation methods have been introduced. These approaches train the student on sequences it generates dynamically, with the teacher providing supervision on these self-generated samples (Agarwal et al., 2024). While this methodology effectively aligns the training and inference distributions, its iterative nature—requiring repeated generation and training steps—incurs substantial computational overhead. To harness the benefits of on-policy training while maintaining offline efficiency, our work introduces a novel offline adaptation of this principle. Our framework is executed in a straightforward two-stage process: first, the student model conducts a single inference pass over the training corpus to generate a complete set of responses. Second, conventional logits-based distillation is performed using this student-generated dataset, with the teacher model providing the supervisory soft labels. This method effectively simulates an on-policy data distribution, thereby alleviating exposure bias while circumventing the costly iterative loop inherent to true on-policy methods. The result is a simple and scalable framework that markedly improves training efficiency over online counterparts without a significant compromise in model performance.

3. Post-training Strategy

In this section, we present the post-training strategy for our base model openPangu Embedded-1B. This strategy includes Two-Stage Curriculum SFT, and Offline On-policy Knowledge Distillation, as shown in Fig 1.

3.1. Post-training Data

For post-training data, core principles of high quality, diversity, and complexity are prioritized, with a reasoning-centered design to enhance the model’s capabilities and generalization as well. The initial data pool integrates multi-source data: open-source instruction datasets, real-world industrial queries (e.g., finance/healthcare scenarios), and synthetic problems derived from pre-training corpora. The post-training data are split into two key subsets: reasoning tasks (advanced STEM with multi-step computation, code generation requiring symbolic manipulation, logical inference) and non-reasoning tasks (general QA, text analysis, long-context understanding, semantic classification, tool use, and agent), with a 3:1 sampling ratio adopted.

To ensure data quality, a rigorous two-stage processing pipeline is constructed: 1) Prior filtering: we leverage

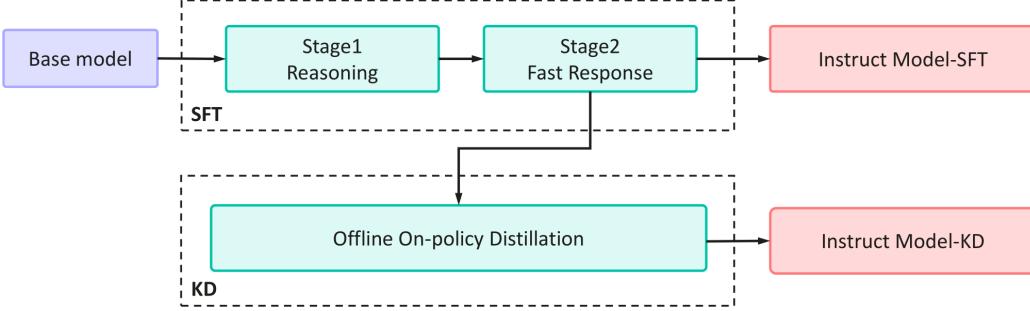


Figure 1. An illustration of the openPangu Embedded-1B post-training pipeline. The pipeline consists of two primary stages: Two-Stage Curriculum SFT and Offline On-policy Knowledge Distillation.

our models to annotate data with attributes (subcategory, question type, answer verifiability, difficulty metrics like reasoning hops) to filter unqualified samples; 2) Diversity maintenance: We use N-gram-based MinHash-LSH to eliminate near-duplicates, followed by the ZIP algorithm (guided by entropy (Yin et al., 2024)) to select samples by prioritizing low compression ratios (higher diversity) and minimizing similarity to existing entries, to yield a pattern-rich dataset. The overall post-training dataset emphasizes reasoning tasks as well as dataset diversity, thus strengthening the model’s ability to avoid superficial pattern matching, while non-reasoning tasks support basic language competence, collectively enhancing performance across both specialized reasoning tasks and general abilities.

3.2. Two-Stage Curriculum SFT

During the SFT phase, we employ Two-Stage Curriculum SFT (Chen et al., 2025), a method grounded in cognitive science: robust reasoning skills must first be explicitly learned before they can be applied intuitively. In Stage 1, the model is trained on reasoning-enhanced data containing explicit, step-by-step chains of thought to build strong inferential capabilities and avoid shallow heuristics. In Stage 2, training shifts to standard prompt–response pairs without intermediate reasoning steps, encouraging the model to implicitly apply its learned reasoning framework to produce concise, accurate outputs. This approach effectively enhances reasoning ability, overall performance, and response efficiency.

3.3. Offline On-policy Knowledge Distillation

On-policy distillation (Agarwal et al., 2024; Lin et al., 2020) has been proven to be an effective model compression technique. It leverages the outputs of a larger, more powerful teacher model to guide a smaller student model, thereby enhancing the student’s accuracy and generalization capabilities. This process enables compact student models to achieve superior performance.

To decouple the generation of teacher model guidance data

from the training process of the student model, and to allow independent optimization of the training data and process, we design an *offline version of on-policy distillation*. Unlike the online version, which cannot dynamically adjust and optimize the quality of the generated logits, the offline version allows for selecting high-quality data during the preprocessing stage to generate teacher model guidance information. Compared to the standard SFT process, our method introduces only an additional Knowledge Distillation (KD) loss term. This design simplifies the implementation process, ensuring ease of use and straightforward operation, while enhancing the flexibility and controllability of the system.

Our method is systematically structured into two distinct phases: an offline data preparation phase followed by a training phase. The offline data preparation phase includes two key components: Student-driven Response Generation and Teacher’s Token-Level Logits Prediction:

- **Student-driven Response Generation:** The process commences with our SFT-trained student model, which performs inference on the queries from the original training dataset. This generates an intermediate on-policy dataset, D_s , where each sample consists of a query and the corresponding student-generated response. This initial step, which we refer to as *distillation by student*, ensures that the subsequent teacher guidance is grounded in the student’s own output distribution.
- **Teacher’s Token-Level Logits Prediction:** Following response generation, we employ a powerful teacher model(e.g., a 7B parameter model sharing the same tokenizer) to annotate each sequence in D_s with token-level guidance. For each position n in a student-generated sequence, the teacher model is conditioned on the prefix of the preceding $n-1$ tokens to predict the distribution for the n -th token. Crucially, rather than performing greedy decoding, we preserve the teacher’s full predictive distribution by recording the logits for the top-k most probable tokens. This constrained con-

ditioning strategy is the cornerstone of our approach. By compelling the teacher to generate guidance from a probabilistic space already accessible to the student, we effectively minimize the intrinsic distributional divergence between the two models. This facilitates a more stable and efficient knowledge transfer.

Training with a Composite Loss Function. In the training phase, the student model is optimized using a composite loss function that synergistically combines standard supervised learning with knowledge distillation. The total loss L_{total} is formulated as the weighted sum of the Cross-Entropy (CE) loss L_{CE} and the knowledge distillation loss L_{KD} :

$$L_{total} = (1 - \lambda_{KD}) \cdot L_{CE} + \lambda_{KD} \cdot L_{KD}$$

where λ_{KD} is a scalar hyperparameter that is a weighting coefficient. It meticulously balances the influence of the direct supervised objective (L_{CE}) and the teacher’s distributional guidance (L_{KD}). The core of the distillation process is the minimization of the KL divergence, denoted $D_{KL}(P \parallel Q)$, is an asymmetric measure that quantifies how a probability distribution Q (from the student) differs from a reference probability distribution P (from the teacher). For discrete distributions over a vocabulary of classes C , it is defined as:

$$D_{KL}(P \parallel Q) = \sum_{c \in C} P(c) \log \frac{P(c)}{Q(c)}$$

By minimizing the KL divergence, the student’s output distribution (Q) is trained to approximate the soft-target distribution (P) provided by the teacher. A lower divergence value signifies that the student has successfully learned to mimic the teacher’s predictive patterns, effectively internalizing the knowledge transferred during distillation.

4. Post-training Evaluation

4.1. Main Results

In this section, we evaluate openPangu Embedded-1B-KD, which is trained using our post-training techniques, on both reasoning and normal language tasks.

SFT Training Setup. The SFT phase of openPangu Embedded-1B uses a two-stage fine-tuning approach, each stage running for 10 epochs. For overall training stability, we employ the AdamW optimizer with a weight decay of 0.1 and apply gradient clipping at a threshold of 1.0. We set the maximum sequence length to 32K to maximize computational efficiency, packing multiple samples into each sequence.

The first stage focuses on complex reasoning, using a global batch size of 4 million tokens. Its learning rate follows

a cosine schedule with a 200-iteration warmup, annealing from a peak of 2×10^{-5} down to 2×10^{-6} . The second stage targets open-ended generation tasks, using a smaller global batch size of 2 million tokens. The learning rate for this stage also follows a cosine schedule, decaying from 1×10^{-5} to 1×10^{-6} .

Baselines & Benchmarks. The openPangu Embedded-1B family is compared against several prominent open-source models within a similar parameter class to ensure a relevant and competitive analysis. The compared baselines include Qwen3 (1.7B and 0.6B) (Yang et al., 2025), Qwen2.5(1.5B) (Yang et al., 2024), Gemma3 (1B) (Team et al., 2024), Llama3.2 (1B) (Dubey et al., 2024) and MiniCPM4 (0.5B) (Team et al., 2025). Including the larger Qwen3-1.7B model serves as a critical point of comparison, allowing for an evaluation of the parameter efficiency of the proposed approach. Performance is measured using standard metrics appropriate for each benchmark. Accuracy (Acc) is used for tasks with single correct answers, such as MMLU (Hendrycks et al., 2020) and GSM8K (Cobbe et al., 2021). The F1 Score is employed for tasks like DROP (Dua et al., 2019), which require a balance of precision and recall in text extraction. For code generation tasks like MBPP (Austin et al., 2021) and HumanEval (Chen et al., 2021), Pass@1 is used, which measures the percentage of problems for which a correct solution is generated in a single attempt.

To ensure a fair and comprehensive evaluation, the evaluation suite is meticulously curated to probe a wide range of cognitive abilities, categorized into four primary domains:

- **General tasks:** Fundamental language understanding, multi-domain knowledge, and Chinese language proficiency are assessed using established benchmarks. These include MMLU for broad, multi-disciplinary knowledge; CMMLU (Li et al., 2023) and C-Eval (Huang et al., 2023) for comprehensive Chinese language evaluation; IF-Eval (Zhou et al., 2023b) for instruction-following fidelity; and CLUEWSC (Xu et al., 2020) for commonsense reasoning. These benchmarks test a model’s core knowledge base and ability to apply it in varied contexts.
- **Mathematics:** To evaluate complex, multi-step quantitative reasoning, the evaluation employs GSM8K and MATH-500 (Hendrycks et al., 2021). These benchmarks require numerical computation and the critical ability to translate natural language problems into logical steps, testing the depth of a model’s reasoning capacity.
- **Reasoning:** The models’ capacity for complex reasoning and information extraction is tested using

Table 1. Instruct model (non-thinking) comparison between openPangu Embedded-1B-KD and other representative models across a diverse set of benchmarks for evaluating language and reasoning skills. **Bold** values represent the best results in each line among models at the 1B-parameter scale. If the original paper reports the results, we present the results from the original paper (marked with asterisks *); otherwise, we list our reproduced results.

Benchmark (Metric)		Qwen3	Qwen2.5	Gemma3	Llama3.2	Qwen3	MiniCPM4	openPangu Embedded SFT	openPangu Embedded KD
# Total Params		1.7B	1.5B	1B	1B	0.6B	0.5B	1B	1B
General	MMLU(Acc)	63.37	52.84	37.49	32.19	44.24	55.55*	63.21	67.28
	CMMLU (Acc)	61.22	53.98	31.57	10.81	42.94	65.22*	53.10	57.75
	C-Eval (Acc)	61.00*	59.30	32.49	32.08	42.60*	66.11*	58.51	66.55
	IF-Eval (Prompt Strict)	68.20*	42.50*	51.57	39.37	54.50*	50.28	56.38	62.66
	CLUEWSC(Acc)	77.36	74.59	50.20	52.36	50.31	49.90	76.95	80.02
Math	GSM8K (Acc)	77.03	73.20*	57.16	36.39	59.29	52.08*	70.89	77.33
	MATH-500 (Acc)	73.00*	46.60	39.40	18.20	55.20*	29.60*	56.20	73.80
Reasoning	DROP (F1)	61.21	48.44	30.98	45.23	34.69	30.07	28.73	40.95
	GPQA-Diamond (Pass@1)	28.60*	29.80*	19.20*	29.29	22.90*	28.28	43.43	44.44
Code	MBPP (Pass@1)	60.70	63.20*	58.75	40.08	46.69	59.14*	52.53	61.09
	HumanEval (Pass@1)	68.90	61.60*	40.24	29.88	40.85	46.34*	59.15	65.85
Average		63.69	55.10	40.82	33.26	44.93	48.42	56.28	63.43

DROP (Dua et al., 2019), which measures reading comprehension intertwined with arithmetic reasoning, and GPQA-Diamond (Rein et al., 2024). This question-answering dataset that probes deep, domain-specific reasoning in physics and chemistry.

- **Code Generation:** Proficiency in programming is evaluated using MBPP (Austin et al., 2021) and HumanEval (Chen et al., 2021). These benchmarks assess the ability to synthesize correct and functional code from natural language docstrings, a key skill for practical applications.

Evaluation Results. The empirical results (summarized in Table 1) unequivocally establish the superior performance of the openPangu Embedded model family, with the openPangu Embedded-1B-KD setting a new state-of-the-art for models in the 1B parameter class. A salient finding is its aggregate performance, achieving an average score of 63.43, which is on par with the larger Qwen3-1.7B model (63.69). This demonstrates exceptional parameter efficiency, suggesting that advanced training and alignment methodologies can be more impactful than simply scaling model size. The model’s most profound advantage lies in mathematical and complex reasoning, where it achieves leading scores on GSM8K (77.33) and MATH-500 (73.80). This strength is complemented by robust general knowledge capabilities, securing top positions on benchmarks like CLUEWSC (80.02) and MMLU (67.28), highlighting a robust and well-rounded bilingual foundation.

4.2. SFT

To optimize the SFT process, we conduct a series of ablation studies to investigate whether a multi-stage training curriculum can enhance the model’s “fast thinking” or intuitive

response capabilities. Specifically, we seek to determine whether first fine-tuning on deliberative reasoning tasks before fine-tuning on rapid-response data yields superior performance. We evaluate three distinct training methods:

- **One-Stage Curriculum (Direct Fast Response).** In this single-stage approach, the model is directly fine-tuned using only the fast response dataset. This baseline measures the efficacy of training exclusively on target-domain data without any preparatory learning phases.
- **Two-Stage Curriculum (Reasoning-to-Fast, without CoT).** This method includes a two-stage curriculum. The model is first trained on a Reasoning dataset from which the intermediate reasoning steps (i.e., the “thought process”) have been explicitly removed. Following this, the model is fine-tuned on the fast response dataset. This approach tests the benefit of a sequential training regimen on datasets with different characteristics, without explicitly teaching the model to reason.
- **Two-Stage Curriculum (Reasoning-to-Fast, with CoT).** openPangu Embedded-1B is also trained in two stages. However, in the initial stage, it is fine-tuned on the complete Reasoning dataset, which crucially retains the detailed reasoning chains and thought processes. This stage is designed to instill deliberative reasoning capabilities before fine-tuning for rapid-response generation.

The experimental outcomes highlight the efficacy of our proposed two-stage curriculum, as both the “Reasoning (w/ CoT)-to-Fast” and “Reasoning (w/o CoT)-to-Fast” approaches significantly outperform the direct “Fast-Response” fine-tuning method. Notably, the two curriculum

Table 2. Accuracy comparison of SFT methods with greedy decoding. The Reasoning (w/ CoT) to Fast approach can effectively improve accuracy during the SFT stage.

Method	General					Math		Reasoning		Code		AVG
	MMLU	CMMLU	C-Eval	IF-Eval	CLUEWSC	GSM8K	MATH-500	DROP	GPQA-Diamond	MBPP	HumanEval	
Direct Fast-Response	60.83	49.58	60.06	56.19	71.21	70.20	46.60	55.84	32.83	54.47	54.88	55.70
Reasoning (w/o CoT) → Fast	61.76	52.02	62.66	61.18	74.18	70.36	51.60	42.60	34.85	54.86	56.71	56.62
Reasoning (w/ CoT) → Fast	63.21	53.10	58.51	56.38	76.95	70.89	56.20	28.73	43.43	52.53	59.15	56.28

variants perform comparably—with the non-CoT approach achieving a peak accuracy of 56.62%—indicating that the curriculum structure itself is the primary driver of success, even without explicit reasoning paths. This initial reasoning phase acts as a form of cognitive scaffolding, equipping the model with foundational problem-solving skills. During the second fast response phase, the model is not merely memorizing input-output pairs but can leverage its acquired reasoning abilities to generate more robust and accurate responses. These findings demonstrate the profound value of a curriculum that prioritizes the development of underlying skills before optimising for rapid task completion.

4.3. Knowledge Distillation

We conduct a systematic ablation study on knowledge distillation to further enhance model performance. Our analysis focuses on four key dimensions: (i) the weighting of the distillation loss term, (ii) the effect of the top-k value during decoding, (iii) the choice of distillation strategy. This investigation aims to elucidate each factor’s relative importance and identify the most effective configuration.

Knowledge Distillation Loss Weight. To optimize the balance between the standard cross-entropy loss and the distillation loss, we conduct experiments with different values of the KD loss weight (λ_{KD}), ranging from 0.5 to 1.0. To accelerate the experimental cycle, we conduct distillation experiments based on a single stage (Direct Fast Response) of SFT. This approach introduces a new Stage 2 to perform distillation guided by labels, with the distillation learning rate kept consistent with that used in SFT. This systematic comparison shows setting $\lambda_{KD} = 0.9$ yields the best overall performance, as shown in Table 3.

Table 3. Effect of KD loss weight on the average benchmark performance of stage2 distillation by label. The evaluation metric is the average zero-shot accuracy across eight benchmarks: MMLU, CMMLU, CEVal, BBH, GSM8K, MATH, MBPP, and HumanEval. All models use greedy decoding with a decode length of 8K.

λ_{KD}	0.5	0.6	0.7	0.8	0.9	1.0
Average	53.94	54.07	53.88	54.43	55.29	54.44

Top-k Value Effect. To explore the impact of the number of top tokens used in the distillation process, we experiment

with four values for the top-k parameter: 5, 10, 15 and 20. By adjusting the value of k, we aim to investigate how the number of candidate logits influences model performance. The experimental setup is consistent with that described in the *Knowledge Distillation Loss Weight* section. We observe that training time remains almost identical across different k values. In terms of performance, increasing k from 5 to 10 yields an improvement of 0.51%, with top-k=10 achieving the highest average accuracy. However, further enlarging k to 15 or 20 leads to a degradation in performance, as shown in Table 4. Hence, we adopt $\lambda_{KD} = 0.9$ and top-k=10 by default unless otherwise specified.

Table 4. Effect of top-k value on the average benchmark performance of stage2 distillation by label. The evaluation metric is the average zero-shot accuracy across eight benchmarks: MMLU, CMMLU, CEVal, BBH, GSM8K, MATH, MBPP, and HumanEval. All models use greedy decoding with a decode length of 8K, and the KD loss weight is fixed at $\lambda_{KD} = 0.9$.

Top-k	5	10	15	20
Average Performance	54.78	55.29	54.24	54.40

Distillation Strategy. We conduct a key ablation study to compare three distinct distillation strategies. To simplify the experimental setup, we introduce an additional Knowledge Distillation (KD) stage to our two-stage Supervised Fine-Tuning (SFT) framework, specifically for distilling the ‘fast-response’ data. The first is the conventional approach, “Distillation by Label” (Hinton et al., 2015; Sanh et al., 2019), where teacher logits are conditioned on ground-truth labels. The second strategy, “Distillation by Teacher” (Kim & Rush, 2016), conditions the teacher’s logits on its own generated response, which can better approximate the teacher model’s intrinsic data distribution. The third strategy is an offline on-policy approach, “Distillation by Student”, where the teacher provides target logits based on the student model’s own generated response.

As shown in Table 5, all methods prove effective. “Distillation by Teacher” achieves a significant accuracy gain of 3.26%. However, the results demonstrate the clear superiority of the “Distillation by Student” strategy, which achieves a more substantial performance increase of 6%. This finding highlights the benefit of aligning the teacher’s guidance with the student’s current output space. By conditioning

Table 5. Ablation study on knowledge distillation. We compare three approaches: "Distillation by Label" where teacher logits are conditioned on ground-truth labels and guided by its top-10 predictions; "Distillation by Teacher" where logits are conditioned on the teacher's own generated response; and "Distillation by Student" where logits are conditioned on the student's predictions to encourage self-consistency. The asterisk (*) denotes that the student model is updated twice with its latest parameters during training for response generation. All results are based on greedy decoding with a generation length of 8K tokens.

Method	General					Math		Reasoning		Code		AVG
	MMLU	CMMLU	C-Eval	IF-Eval	CLUEWSC	GSM8K	MATH-500	DROP	GPQA-Diamond	MBPP	HumanEval	
Stage2 SFT	63.21	53.10	58.51	56.38	76.95	70.89	56.20	28.73	43.43	52.53	59.15	56.28
Stage3 Distillation by Label	62.68	53.97	58.83	59.70	77.56	72.10	65.20	30.24	46.46	57.20	62.80	58.79
Stage3 Distillation by Teacher	61.74	54.94	62.83	59.89	79.92	74.30	61.20	36.59	47.98	55.25	60.37	59.55
Stage3 Distillation by Student	65.91	56.17	64.30	59.33	79.51	76.72	71.40	40.99	44.95	58.75	67.07	62.28
Stage3 Distillation by Student*	67.28	57.75	66.55	62.66	80.02	77.33	73.80	40.95	44.44	61.09	65.85	63.43

Table 6. Ablation study on post-train pipeline. KD use distillation by label which teacher logits are conditioned on ground-truth labels, with the teacher's topk-10 predicted logits serving as auxiliary guidance. All models use greedy decoding with a decode length of 8K.

Method	General					Math		Reasoning		Code		AVG
	MMLU	CMMLU	C-Eval	IF-Eval	CLUEWSC	GSM8K	MATH-500	DROP	GPQA-Diamond	MBPP	HumanEval	
SFT(Reasoning) + SFT(Fast)	63.21	53.10	58.51	56.38	76.95	70.89	56.20	28.73	43.43	52.53	59.15	56.28
SFT(Reasoning) + KD(Fast)	61.38	55.12	60.14	56.93	75.41	72.78	64.60	43.01	42.93	51.36	61.59	58.66
SFT(Reasoning) + SFT(Fast) + KD(Fast)	62.68	53.97	58.83	59.70	77.56	72.10	65.20	30.24	46.46	57.20	62.80	58.79
SFT(Reasoning) + KD(Fast) + SFT(Fast)	63.75	54.78	60.63	58.78	73.05	66.94	55.20	45.99	43.94	55.25	58.54	57.90
SFT(Reasoning) + SFT(Fast) + KD(Reasoning) + KD(Fast)	62.69	53.95	62.24	57.86	79.61	73.09	66.20	33.61	39.39	56.42	60.37	58.68
SFT(Reasoning) + KD(Reasoning) + SFT(Fast) + KD(Fast)	62.82	53.22	60.56	59.15	81.25	73.16	64.80	37.27	40.40	57.59	65.24	59.59

distillation on the student's response, the teacher provides corrective and refining signals on a distribution immediately relevant to the student's state. This alignment minimizes the distributional mismatch between the two models, creating a more stable learning signal and facilitating more effective knowledge transfer.

Furthermore, we find that the performance of the "Distillation by Student" strategy can be enhanced. By periodically updating the student model with its latest parameters during the training process—in our case, twice—and re-generating responses for subsequent distillation, the model's accuracy is further improved to 63.43%. This creates a dynamic, self-correcting loop that maximizes the efficiency of knowledge transfer and significantly boosts the final performance of the student model.

4.4. Ablation Study on Post-training Pipeline

We further conduct an ablation study to investigate the optimal design of the post-training pipeline, focusing on how to interleave supervised fine-tuning and knowledge distillation. Here, KD is implemented as label-based distillation, where the teacher logits are conditioned on the ground-truth labels, and the teacher's top- $k=10$ predicted logits serve as auxiliary guidance. The results are summarized in Table 6. When comparing the two-stage settings, we find that replacing the final *SFT (Fast)* phase with *KD (Fast)* improves performance (AVG 58.66 vs. 56.28), indicating that knowledge distillation provides stronger supervision than repeated SFT alone. Extending to three-stage pipelines, introducing KD after the initial reasoning-oriented SFT further improves generalization, with the configuration *SFT (Reasoning) +*

SFT (Fast) + KD (Fast) achieving an average score of 58.79, slightly better than its variant with reversed order (57.90). Finally, we explore four-stage pipelines that combine both reasoning and fast-thinking KD. Among these, the sequence *SFT (Reasoning) + KD (Reasoning) + SFT (Fast) + KD (Fast)* achieves the best overall performance (AVG 59.59), delivering strong gains on benchmarks such as CLUEWSC (81.25), GSM8K (73.16), and HumanEval (65.24). This suggests that a balanced alternation of SFT and KD, especially when incorporating both reasoning and fast-thinking teacher guidance, yields the most effective transfer of knowledge in the post-training phase and sets a clear direction for optimizing future training pipelines.

5. Conclusion and Discussion

This work introduces a comprehensive post-training pipeline designed to significantly enhance the capabilities of small language models, combining curriculum-based SFT and offline on-policy knowledge distillation. Applying this post-training pipeline, we have developed a compact yet powerful model that achieves state-of-the-art performance among models with around 1 billion parameters. All the training and deployment are conducted on Ascend hardware. Extensive experiments demonstrate that openPangu Embedded-1B-KD excels in mathematical reasoning, code generation, and multilingual understanding while maintaining superior inference efficiency on Ascend edge hardware. This provides a practical and scalable solution for bridging high-performance AI and edge deployment, paving the way for advanced on-device intelligence. Future work will focus on expanding capabilities and optimizing for broader edge

scenarios.

6. Acknowledgments

The success of this research is a collaborative effort. We are particularly grateful to the Data team for curating the training dataset, the Infrastructure team for building the robust training and inference framework, and the Evaluation team for their insightful assessment of the model’s performance.

References

- Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Agarwal, R., Vieillard, N., Zhou, Y., Stanczyk, P., Garea, S. R., Geist, M., and Bachem, O. On-policy distillation of language models: Learning from self-generated mistakes. In *The twelfth international conference on learning representations*, 2024.
- Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Bai, J., Bai, S., Chu, Y., Cui, Z., Dang, K., Deng, X., Fan, Y., Ge, W., Han, Y., Huang, F., et al. Qwen technical report. *arXiv preprint arXiv:2309.16609*, 2023.
- Chen, H., Wang, Y., Han, K., Li, D., Li, L., Bi, Z., Li, J., Wang, H., Mi, F., Zhu, M., et al. Pangu embedded: An efficient dual-system llm reasoner with metacognition. *arXiv preprint arXiv:2505.22375*, 2025.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., Pondé, H., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., Ryder, N., Pavlov, M., Power, A., Kaiser, L., Bavarian, M., Winter, C., Tillet, P., Such, F. P., Cummings, D. W., Plappert, M., Chantzis, F., Barnes, E., Herbert-Voss, A., Guss, W. H., Nichol, A., Babuschkin, I., Balaji, S., Jain, S., Carr, A., Leike, J., Achiam, J., Misra, V., Morikawa, E., Radford, A., Knight, M. M., Brundage, M., Murati, M., Mayer, K., Welinder, P., McGrew, B., Amodei, D., McCandlish, S., Sutskever, I., and Zaremba, W. Evaluating large language models trained on code. *ArXiv*, abs/2107.03374, 2021. URL <https://api.semanticscholar.org/CorpusID:235755472>.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Du, Q., Zong, C., and Zhang, J. Mods: Model-oriented data selection for instruction tuning. *arXiv preprint arXiv:2311.15653*, 2023.
- Dua, D., Wang, Y., Dasigi, P., Stanovsky, G., Singh, S., and Gardner, M. Drop: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *North American Chapter of the Association for Computational Linguistics*, 2019. URL <https://api.semanticscholar.org/CorpusID:67855846>.
- Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Fu, W., Li, S., Zhao, Y., Ma, H., Dutta, R., Zhang, X., Yang, K., Jin, Y., and Guo, X. Hardware phi-1.5 b: A large language model encodes hardware domain specific knowledge. In *2024 29th Asia and South Pacific Design Automation Conference (ASP-DAC)*, pp. 349–354. IEEE, 2024.
- Groeneveld, D., Beltagy, I., Walsh, P., Bhagia, A., Kinney, R., Tafjord, O., Jha, A. H., Ivison, H., Magnusson, I., Wang, Y., et al. Olmo: Accelerating the science of language models. *arXiv preprint arXiv:2402.00838*, 2024.
- Guo, D., Zhu, Q., Yang, D., Xie, Z., Dong, K., Zhang, W., Chen, G., Bi, X., Wu, Y., Li, Y., et al. Deepseek-coder: When the large language model meets programming—the rise of code intelligence. *arXiv preprint arXiv:2401.14196*, 2024.
- Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., and Steinhardt, J. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020.
- Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., and Steinhardt, J. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.
- Hinton, G., Vinyals, O., and Dean, J. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Huang, Y., Bai, Y., Zhu, Z., Zhang, J., Zhang, J., Su, T., Liu, J., Lv, C., Zhang, Y., Lei, J., Qi, F., Fu, Y., Sun, M., and He, J. C-eval: A multi-level multi-discipline chinese evaluation suite for foundation models. *ArXiv*, abs/2305.08322, 2023. URL <https://api.semanticscholar.org/CorpusID:258685666>.
- Kashyap, R., Kashyap, V., et al. Gpt-neo for commonsense reasoning—a theoretical and practical lens. *arXiv preprint arXiv:2211.15593*, 2022.

- Kim, Y. and Rush, A. M. Sequence-level knowledge distillation. In *Proceedings of the 2016 conference on empirical methods in natural language processing*, pp. 1317–1327, 2016.
- Li, H., Zhang, Y., Koto, F., Yang, Y., Zhao, H., Gong, Y., Duan, N., and Baldwin, T. Cmmlu: Measuring massive multitask language understanding in chinese. *arXiv preprint arXiv:2306.09212*, 2023.
- Li, M., Zhang, Y., Li, Z., Chen, J., Chen, L., Cheng, N., Wang, J., Zhou, T., and Xiao, J. From quantity to quality: Boosting LLM performance with self-guided data selection for instruction tuning. In Duh, K., Gomez, H., and Bethard, S. (eds.), *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 7602–7635, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.naacl-long.421. URL <https://aclanthology.org/2024.naacl-long.421/>.
- Lin, A., Wohlwend, J., Chen, H., and Lei, T. Autoregressive knowledge distillation through imitation learning. *arXiv preprint arXiv:2009.07253*, 2020.
- Liu, A., Feng, B., Xue, B., Wang, B., Wu, B., Lu, C., Zhao, C., Deng, C., Zhang, C., Ruan, C., et al. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*, 2024a.
- Liu, W., Zeng, W., He, K., Jiang, Y., and He, J. What makes good data for alignment? a comprehensive study of automatic data selection in instruction tuning. In *The Twelfth International Conference on Learning Representations*.
- Liu, Z., Zhao, C., Iandola, F., Lai, C., Tian, Y., Fedorov, I., Xiong, Y., Chang, E., Shi, Y., Krishnamoorthi, R., et al. Mobilellm: Optimizing sub-billion parameter language models for on-device use cases. In *Forty-first International Conference on Machine Learning*, 2024b.
- Lobo, E., Agarwal, C., and Lakkaraju, H. On the impact of fine-tuning on chain-of-thought reasoning. *arXiv preprint arXiv:2411.15382*, 2024.
- Luong, T. Q., Zhang, X., Jie, Z., Sun, P., Jin, X., and Li, H. Reft: Reasoning with reinforced fine-tuning. *arXiv preprint arXiv:2401.08967*, 2024.
- Rein, D., Hou, B. L., Stickland, A. C., Petty, J., Pang, R. Y., Dirani, J., Michael, J., and Bowman, S. R. Gpqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*, 2024.
- Ren, M., Cao, B., Lin, H., Liu, C., Han, X., Zeng, K., Wan, G., Cai, X., and Sun, L. Learning or Self-aligning? Rethinking Instruction Fine-tuning, August 2024. URL <http://arxiv.org/abs/2402.18243>. arXiv:2402.18243 [cs].
- Ross, S., Gordon, G., and Bagnell, D. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the fourteenth international conference on artificial intelligence and statistics*, pp. 627–635. JMLR Workshop and Conference Proceedings, 2011.
- Sanh, V., Debut, L., Chaumond, J., and Wolf, T. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Tang, Y., Yin, Y., Wang, Y., Zhou, H., Pan, Y., Guo, W., Zhang, Z., Rang, M., Liu, F., Zhang, N., et al. Pangu ultra moe: How to train your big moe on ascend npus. *arXiv preprint arXiv:2505.04519*, 2025.
- Team, G., Mesnard, T., Hardin, C., Dadashi, R., Bhupatiraju, S., Pathak, S., Sifre, L., Rivière, M., Kale, M. S., Love, J., et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024.
- Team, M., Xiao, C., Li, Y., Han, X., Bai, Y., Cai, J., Chen, H., Chen, W., Cong, X., Cui, G., et al. Minicpm4: Ultra-efficient llms on end devices. *arXiv preprint arXiv:2506.07900*, 2025.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Xu, L., Hu, H., Zhang, X., Li, L., Cao, C., Li, Y., Xu, Y., Sun, K., Yu, D., Yu, C., et al. Clue: A chinese language understanding evaluation benchmark. *arXiv preprint arXiv:2004.05986*, 2020.
- Yang, A., Li, A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Gao, C., Huang, C., Lv, C., et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Yang, Q. A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Li, C., Liu, D., Huang, F., Dong, G., Wei, H., Lin, H., Yang, J., Tu, J., Zhang, J., Yang, J., Yang, J., Zhou, J., Lin, J., Dang, K., Lu, K., Bao, K., Yang, K., Yu, L., Li, M., Xue, M., Zhang, P., Zhu, Q., Men, R., Lin, R., Li, T., Xia, T., Ren, X., Ren, X., Fan, Y., Su, Y., Zhang, Y.-C., Wan, Y., Liu, Y., Cui, Z., Zhang, Z., Qiu, Z., Quan, S., and Wang, Z. Qwen2.5 technical report. *ArXiv*, abs/2412.15115, 2024. URL <https://api.semanticscholar.org/CorpusID:274859421>.

Yin, M., Wu, C., Wang, Y., Wang, H., Guo, W., Wang, Y., Liu, Y., Tang, R., Lian, D., and Chen, E. Entropy law: The story behind data compression and lilm performance, 2024. URL <https://arxiv.org/abs/2407.06645>.

Zhang, P., Zeng, G., Wang, T., and Lu, W. Tinyllama: An open-source small language model. *arXiv preprint arXiv:2401.02385*, 2024.

Zhang, S., Roller, S., Goyal, N., Artetxe, M., Chen, M., Chen, S., Dewan, C., Diab, M., Li, X., Lin, X. V., et al. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022.

Zhou, C., Liu, P., Xu, P., Iyer, S., Sun, J., Mao, Y., Ma, X., Efrat, A., Yu, P., Yu, L., Zhang, S., Ghosh, G., Lewis, M., Zettlemoyer, L., and Levy, O. LIMA: Less Is More for Alignment, May 2023a. URL <http://arxiv.org/abs/2305.11206>. arXiv:2305.11206 [cs].

Zhou, J., Lu, T., Mishra, S., Brahma, S., Basu, S., Luan, Y., Zhou, D., and Hou, L. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023b.